

Tidy data

Table of contents

tidyverse	2
Introduction to the tidyverse	3
Dataset	4
Exercise 1.1	5
Piping	6
Exercise 1.2	12
Exercise 1.3	12
Exercise 1.4	13
Exercise 1.5	13
Exercise 1.6	13
Exercise 1.7	13
Exercise 1.8	13
Exercise 1.9	14
Exercise 1.10	14

Copy this code and paste it into a new quarto document in R. We will be using quarto for assignments, it is good to practice using it with each workshop.

tidyverse

By the phrase `tidy data`, it is meant the preferred way of arranging data that is easy to analyse. The principles of tidy data are:

- Each variable forms a column.
- Each observation forms a row.
- Each type of observational unit forms a table.

Introduction to the tidyverse

We will be largely using the `tidyverse` suite of packages for data organisation, summarizing, and plotting; see <https://www.tidyverse.org/>.

Let's load that package now: Remember if you have not installed it you will need to use `install.packages()` first. We installed it in the first workshop so you can return to that code to see how we did it.

```
library(tidyverse)
```

Dataset

For this workshop we will use some tidyverse built in datasets. Each dataset below shows the same values of four variables: country, year, population, and number of documented cases of TB (tuberculosis), but each dataset organizes the values in a different way. Take a look at these datasets by typing their names into a code chunk or directly into the console. You can also try your hand at the functions `head()` and `summary()`.

```
table1
```

```
# A tibble: 6 x 4
  country     year   cases population
  <chr>      <dbl>   <dbl>       <dbl>
1 Afghanistan 1999    745  19987071
2 Afghanistan 2000   2666  20595360
3 Brazil      1999  37737 172006362
4 Brazil      2000  80488 174504898
5 China       1999 212258 1272915272
6 China       2000 213766 1280428583
```

```
table2
```

```
# A tibble: 12 x 4
  country     year type        count
  <chr>      <dbl> <chr>      <dbl>
1 Afghanistan 1999 cases      745
2 Afghanistan 1999 population 19987071
3 Afghanistan 2000 cases     2666
4 Afghanistan 2000 population 20595360
5 Brazil      1999 cases     37737
6 Brazil      1999 population 172006362
7 Brazil      2000 cases     80488
8 Brazil      2000 population 174504898
9 China       1999 cases     212258
10 China      1999 population 1272915272
```

```
11 China      2000 cases      213766
12 China      2000 population 1280428583
```

table3

```
# A tibble: 6 x 3
  country     year   rate
  <chr>       <dbl> <chr>
1 Afghanistan 1999  745/19987071
2 Afghanistan 2000  2666/20595360
3 Brazil      1999  37737/172006362
4 Brazil      2000  80488/174504898
5 China       1999  212258/1272915272
6 China       2000  213766/1280428583
```

Exercise 1.1

For each of the sample tables, describe what each observation and each column represents.
Which is the most tidy?

Piping



Tip

The piping operation is a fundamental aspect of computer programming. The semantics of pipes is taking the output from the left-hand side and passing it as input to the right-hand side.

The R package `magrittr` introduced the pipe operator `%>%` and can be pronounced as “then”. In RStudio windows/Linux versions, press **Ctrl+Shift+M** to insert the pipe operator. On a Mac, use **Cmd+Shift+M**.

R also has its own pipe, `|>`, which is an alternative to `%>%`. You will see both used in this course. If you want to change the pipe inserted automatically with **Ctrl+Shift+M**, find on the menu **Tools > Global Options**, then click on **Code** and check the box that says “**Use Native Pipe Operator**”.

Consider the study guide dataset `rangitikei.csv` (Recreational Use of the Rangitikei river). The first 10 rows of this dataset are shown below:

	<code>id</code>	<code>loc</code>	<code>time</code>	<code>w.e</code>	<code>cl</code>	<code>wind</code>	<code>temp</code>	<code>river</code>	<code>people</code>	<code>vehicle</code>
1	1	1	2	1	1	2	2	1	37	15
2	2	1	1	1	1	2	1	2	23	6
3	3	1	2	1	1	2	2	3	87	31
4	4	2	2	1	1	2	1	1	86	27
5	5	2	1	1	1	2	2	2	19	2
6	6	2	2	1	2	1	3	3	136	23
7	7	1	2	2	2	2	2	3	14	8
8	8	1	2	1	2	2	2	3	67	26
9	9	1	1	2	1	3	1	2	4	3
10	10	2	2	1	2	2	2	3	127	45

Try the following examples after loading the `rangitikei` dataset.

```
select()
```

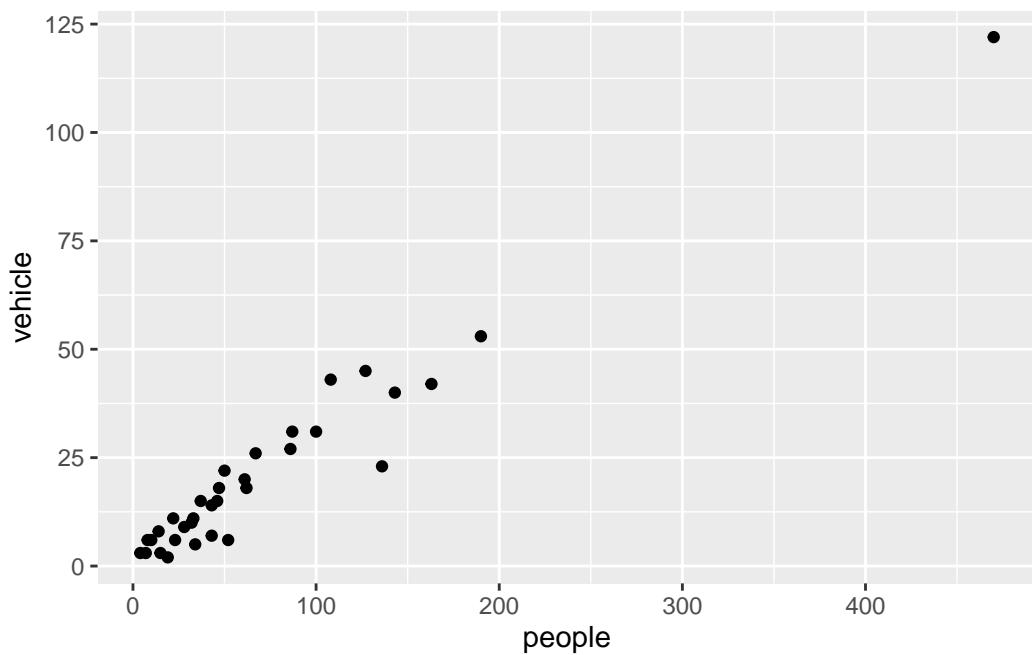
```
new.data <- rangitikei |>  
  select(people, vehicle)
```

```
names(new.data)
```

```
[1] "people"  "vehicle"
```

What does `select()` do?

```
rangitikei |>  
  select(people, vehicle) |> # select columns  
  ggplot() + # make a plot using those columns  
  aes(x=people, y=vehicle) +  
  geom_point()
```

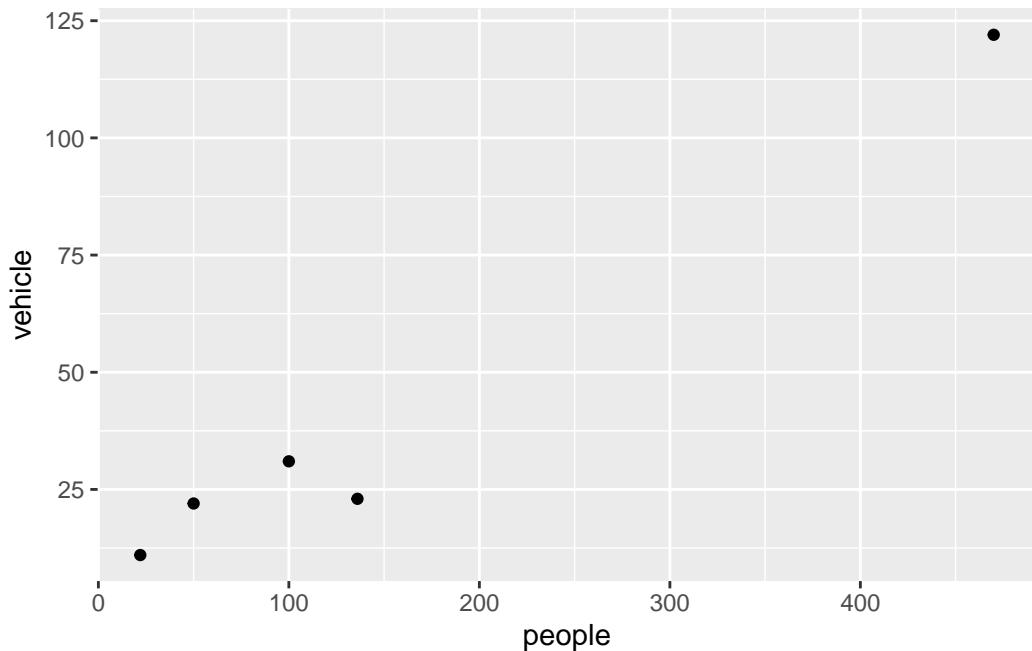


We select two columns and create a scatter plot with the above commands.

Now try another function:

```
filter()
```

```
rangitikei |>
  filter(wind==1) |>
  select(people, vehicle) |>
  ggplot() +
  aes(x=people, y=vehicle) +
  geom_point()
```



What does `filter()` do?

The above commands filter the data for the low wind days and plots vehicle against people. `filter()` subsets the data for all observations matching a specified criteria.

`arrange()`

```
rangitikei |>
  filter(wind==1) |>
  arrange(w.e) |>
  select(w.e, people, vehicle)
```

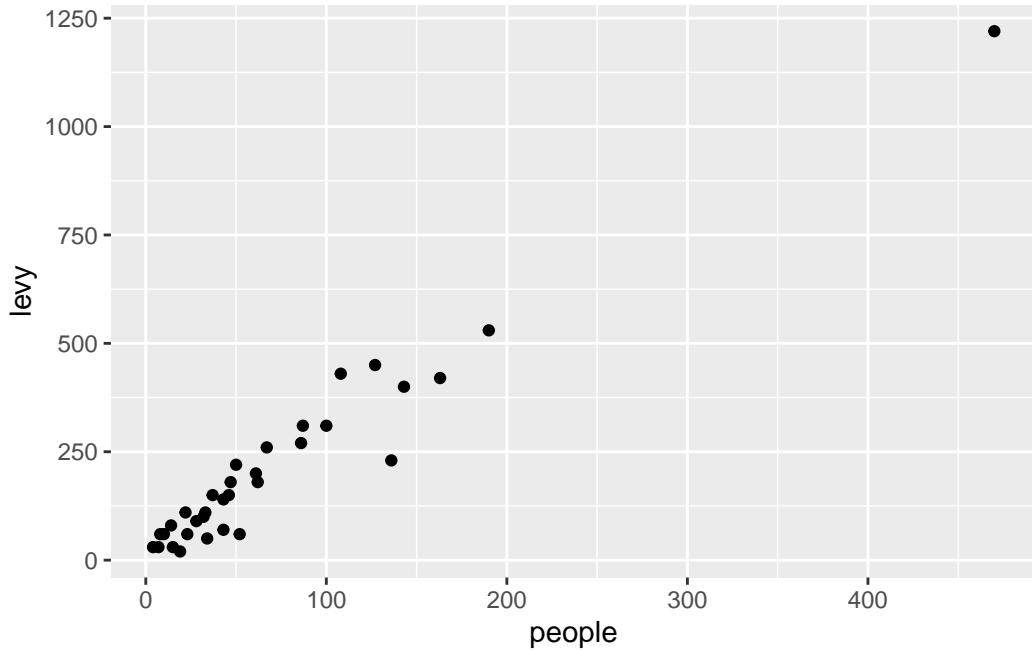
w.e	people	vehicle
1	136	23
2	50	22
3	100	31

```
4     1      470      122
5     2       22       11
```

```
mutate()
```

Assume that a \$10 levy is collected for each vehicle. We can create this new `levy` column as follows.

```
rangitikei |>
  mutate(levy = vehicle*10) |>
  select(people, levy) |>
  ggplot() +
  aes(x = people, y=levy) +
  geom_point()
```



Note that the pipe operation was used to create a scatter plot using the newly created column.

```
summarise()
```

```
rangitikei |>
  summarise(total = n(),
            avg = mean(people))
```

```
)
```

```
total      avg
1     33 71.72727
```

We obtain the selected summary measures namely the total and the mean number of people.
Try-

```
rangitikei |>
  filter(wind == 1) |>
  summarise(total = n(),
            avg = mean(people)
  )
```

```
total      avg
1     5 155.6
```

`group_by()`

We obtain the wind group-wise summaries below:

```
rangitikei |>
  group_by(wind) |>
  summarise(total=n(),
            avg=mean(people))

# A tibble: 3 x 3
  wind total      avg
  <int> <int>    <dbl>
1     1     5 156.
2     2    26  59.7
3     3     2  19
```

There are many more commands such as the `transmute` function which conserves only the needed columns. Try

```
rangitikei |>
  group_by(wind, w.e) |>
  transmute(total=n(),
```

```

    avg=mean(people))

# A tibble: 33 x 4
# Groups:   wind, w.e [6]
  wind   w.e total   avg
  <int> <int> <int> <dbl>
1     2     1     18  72.1
2     2     1     18  72.1
3     2     1     18  72.1
4     2     1     18  72.1
5     2     1     18  72.1
6     1     1      4 189
7     2     2      8 31.8
8     2     1     18  72.1
9     3     2      1  4
10    2     1     18  72.1
# i 23 more rows

```

A simple frequency table is found using `count()`. Try-

```

rangitikei |>
  group_by(wind, w.e) |>
  count(temp)

# A tibble: 10 x 4
# Groups:   wind, w.e [6]
  wind   w.e temp     n
  <int> <int> <int> <int>
1     1     1     1     1
2     1     1     3     3
3     1     2     3     1
4     2     1     1     4
5     2     1     2    12
6     2     1     3     2
7     2     2     2     6
8     2     2     3     2
9     3     1     2     1
10    3     2     1     1

```

```

rangitikei |>
  group_by(wind, w.e) |>
    count(temp, river)

# A tibble: 16 x 5
# Groups:   wind, w.e [6]
  wind   w.e   temp river     n
  <int> <int> <int> <int> <int>
1     1     1     1     1     1
2     1     1     3     3     3
3     1     2     3     3     1
4     2     1     1     1     1
5     2     1     1     2     1
6     2     1     1     3     2
7     2     1     2     1     3
8     2     1     2     2     2
9     2     1     2     3     7
10    2     1     3     3     2
11    2     2     2     1     2
12    2     2     2     3     4
13    2     2     3     2     1
14    2     2     3     3     1
15    3     1     2     2     1
16    3     2     1     2     1

```

The `count()` is useful to check the balanced nature of the data when many subgroups are involved.

Exercise 1.2

Generate a summary table to show the average number of people observed at each of the three rivers.

```
# your code goes here
```

Exercise 1.3

Generate a summary table to show the average number of vehicles observed at each of the three wind levels.

```
# your code goes here
```

Exercise 1.4

Repeat the above exercise 1.2 but now calculate the median number of people and the median number of vehicles at each river in a single table.

```
# your code goes here
```

Exercise 1.5

Use the above three tables you generate to describe the Rangitikei dataset. What other information or tables might be useful?

Now let's practice using these functions using the TB data.

Exercise 1.6

Using `table1`, compute rate of TB cases per 10,000 and the total cases per year

```
# your code goes here
```

Exercise 1.7

Generate a summary table to show the total cases in each country. Your table should have 3 rows.

```
# your code goes here
```

Exercise 1.8

For `table2`, write pseudo-code for how you would perform the following actions. Sketch/describe how you would do these.

- Extract the number of TB cases per country per year.

- b) Extract the matching population per country per year.
- c) Divide cases by population, and multiply by 10000.
- d) Store back in the appropriate place.

Exercise 1.9

What does `na.rm = TRUE` do in the functions `mean()` and `sum()`?

```
# your code goes here
```

Exercise 1.10

What does the function `na.omit()` do? Why would you add it to a piping sequence of code when tidying data?

```
# your code goes here
```